

Notification Systems

System Design

DSA optimization

Code / DB Schema / G-R

Low level Design

microservices / data flow / fault tolerance /
scaling challenges

High level Design

No Code today

5 step approach to any design problem

- ① Problem Statement
 - ② Functional Requirements
 - ③ Non Functional Requirements
 - ④ Scale Estimation
 - ⑤ Design
- Prerequisites

Problem

Notification System - generic notifications

sms
whatsapp
email
app

Scalor

Toy Store

Education Company

E-commerce Company

notification is NOT their core business

core business - something else

notifications - just a feature

Notification.com

SaaS

core offering → Notification SaaS

tons of clients → clients have users

v. high user count

v. high volume

lots of features

Functional Requirements

MVP features

Minimal Viable Product

Notification :

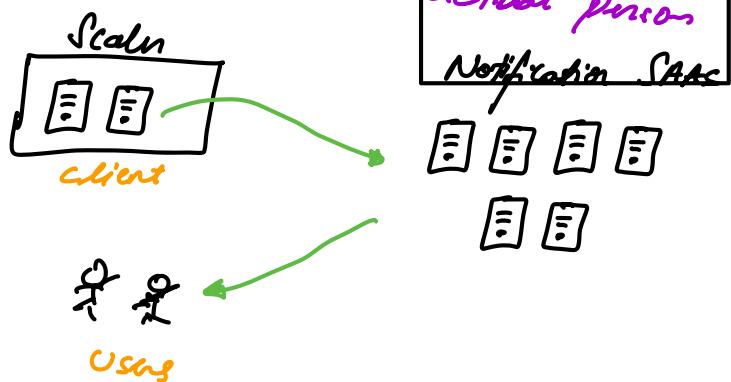
any small content that we wish to send to a user via some means

text/audio/video/images
URL

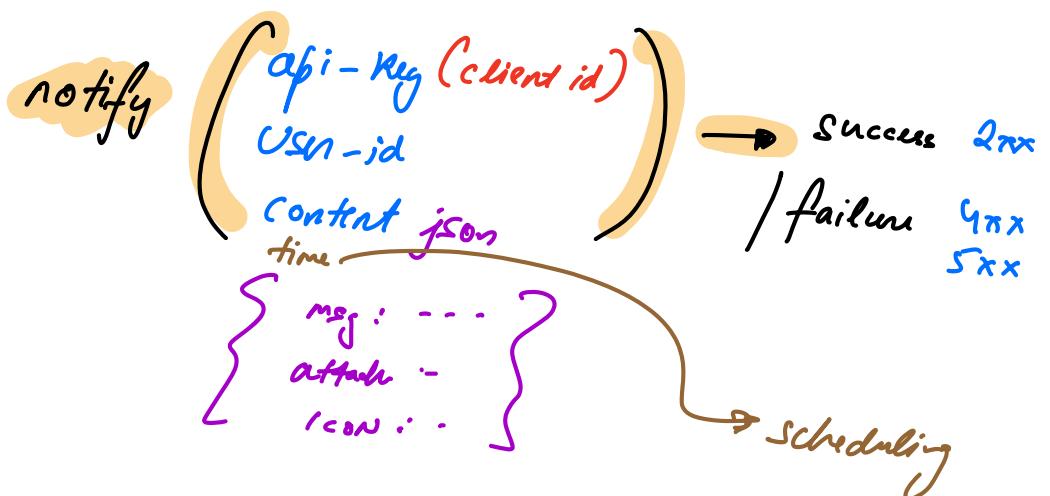
sms/whatsapp/app..
flexible

①

Client can send notification to company's backend



Post



②

System should support multiple channels

whatsapp / sms / email / in-app ...

MUST

be pluggable / extensible
+ slack + call

seamlessly add
more channels
without changing
the architecture

③

Clients should be able to BROADCAST notifications
to multiple users

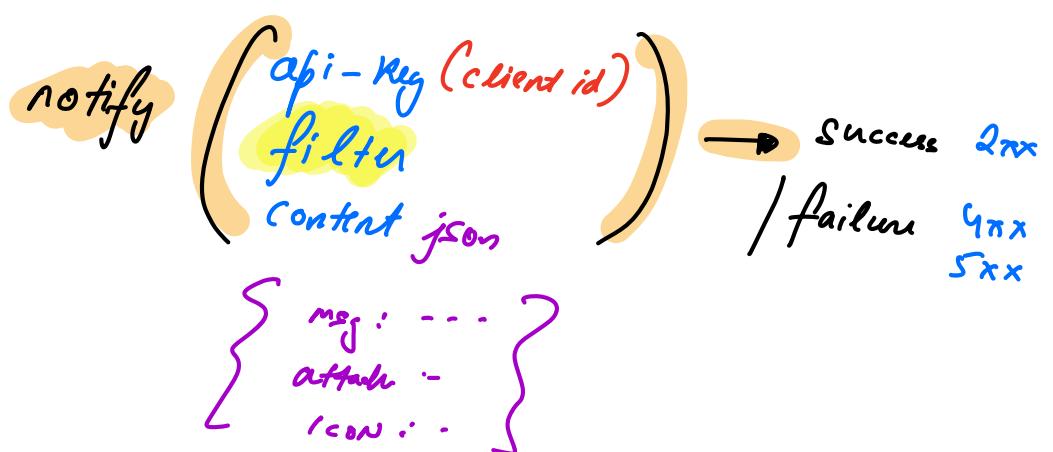
Scalr → Broadcast
Facebook
Reminders

1000 users

Twitter → Elon Musk
Tweets

100 M users

Post

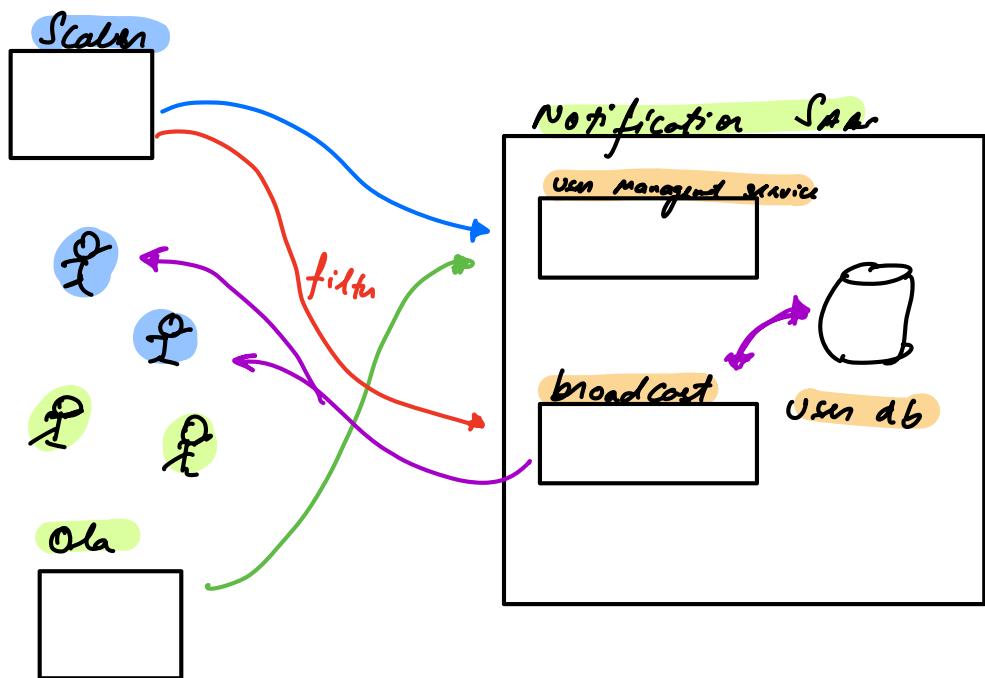


Note: Providing list of user-ids in API call is bad design

4

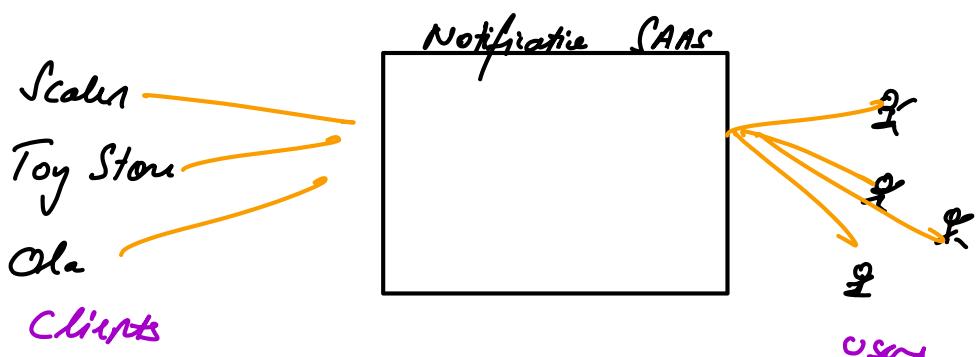
Hidden
Feature

Clients should be able to manage their user list



5

Rate Limiting



① Client limits (Quota limits)

based on the purchase type

free → 500 notifications / day

basic → 50,000

team → 50,000 + additional team

Enterprise → Unlimited

② User side limits

- no user should receive too many notifications

- Users should be able to customize their notification subscriptions

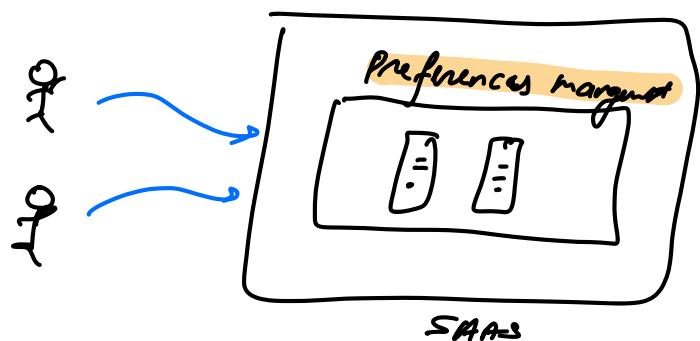
sensitive/critical

OTP / emergency alerts

marketing

channel
— email
— fax
— in-app

Push-notifications



⑥ Observability

- detect something is wrong
- diagnose the issue
- fix
- verify the fix is working

Observability

2020 & 10M
Steve
@google

- # notifications sent / client / day
/ type / channel
- # exceptions / hour

- Logs
- Search & analytics

Non-functional Requirements - Design Goals

① Compliance

TRAI

- very strict rules on what SMS/phone calls we can send
- email providers (SMTP) → own rules

② Priority should be considered

OTPs

v. high priority

notify(. . .)

Promotions

lowest priority



③ High Availability fault tolerant

④ Reliable

backups / retry mechanism

⑤ Low latency

⑥ Single-delivery / Idempotency

∴ can have an auto-retry mechanism

⑦ Security

v. difficult to do correctly
leave it to a security expert

Scale Estimation

100 users only - no need of HLD
Scale matrix

Total # users $\approx 2B$ $(B = 10^9)$
(across all clients)

avg notifications / user / day = 10

Avg # notifications /sec = 2×10^9 user * $\frac{10 \text{ notifications}}{\text{user} * \text{day}}$

$200,000/\text{s}$

$$= 2 \times 10^9 \times 10 \frac{\text{notifications}}{\text{day}}$$

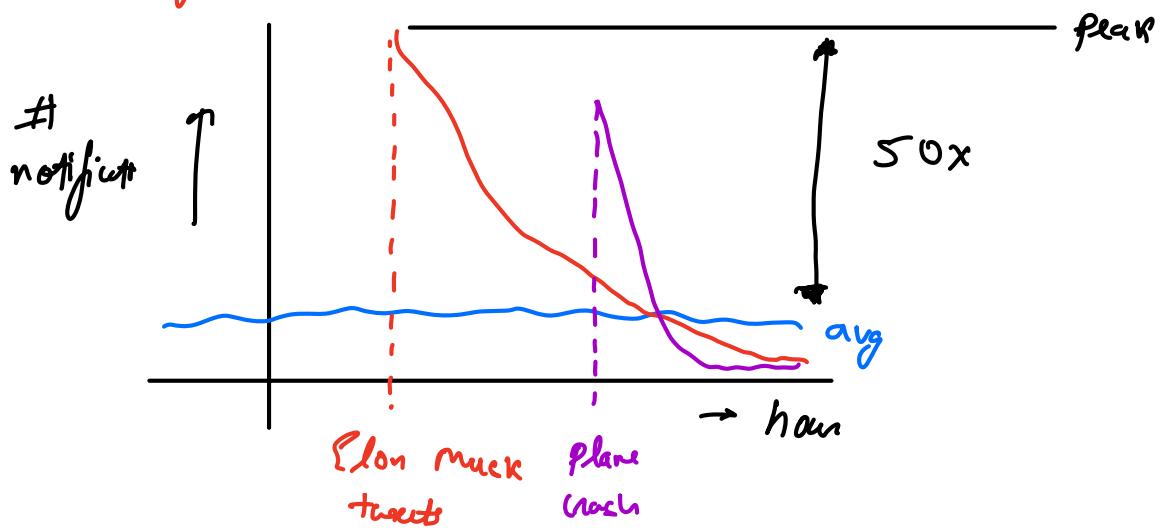
day = $24 \times 60 \times 60$ seconds
 $= 86,400$ seconds
 $\approx 10^5$ seconds

$$= \frac{2 \times 10^9 \times 10}{10^5} \frac{\text{notifications}}{\text{second}}$$

$$= 2 \times 10^5 \text{ notifications/sec}$$

Peak

$$= 10^6 \text{ notifications/sec}$$



$$2 \times 10^5 \text{ avg notification/sec} * 50$$

$$= 10^6 \text{ notifications/sec} \rightarrow \text{Our system should be able to handle this!}$$

System Design



Scalable HLD Curriculum

- How backend systems work
 - DNS / vertical vs horizontal scaling / load balancing / sharding
- Load Balancing & Consistent Hashing
 - Healthcheck & Heartbeat / Roll over / DNS as load balancer / Untrusted Layer / Consistent hashing - deduplication
- Caching
 - Browser / CDN / Backend
 - Local vs Global / Single vs Distributed Redis
 - Cache Invalidation
 - Write through
 - TTL
 - Write back
 - Write around



- **Caching Case Studies**

Scalar Code Judge / Scalar Leaderboard /
Facebook News Feed

- **CAP theorem & Replication**

Availability / Consistency / Partition Tolerance
 /
 Immediate Eventual

CAP theorem / PACELC /

Master-Slave Eventual
 Tunable Consistency Quorum
 / with Everywhere

- **SQL vs NoSQL**

Pros & Cons

Sharding Key

NoSQL Types

Key-Value	graph
Document	Vector
Wide Column	In-memory
Search	timeseries
File store	

- **Data Base Instances**

orchestrator / Config Management / Resend servers /
Shard Creation / Multi-master

- Database Internals
LSM Trees / Bloom Filters / Sparse Index /
Compaction / MemTable / WAL
- Case Study — Google Typahead
Aries / Redis / Recency / Sampling / Ratcheting / Percolation
- Case Study — Messaging apps
Slack groups / WhatsApp / Facebook messenger
- Case Study — Kafka + Zookeeper
event driven architecture / Pub-Sub /
topics & partitions / msg order
- Case Study — Elasticsearch
Inverted Index / TF-IDF / Sharding / Tunable Latency
- Case Study — File Storage
upload / download / chunking
- Case Study — Uber
Quad Tree / traffic aware /

traffic patterns / request driven

- Case Study - Rate Limiting
DDoS / Token Bucket / Leaky Bucket / Sliding Window

- Case Study - Id generation
idempotency / Unique / incremental

- Case Study - Video Streaming
Netflix / YouTube / Riotstar
Concurrency
ABS / chunking / CDNs

- Microservice
 - Monolith vs Microservices vs Service Oriented
 - Breaking a Monolith
 - Observability / Distributed Tracing / Logging
 - Threading herd / Cascading failure / Circuit Breaker
 - Distributed Transactions
 - Two Phase Commit / Saga / CQRS /
 - Orchestration & Choreography
 - Event Sourcing

SDE-3 @ google

60 LPA ... 1-2 Cr base
+ Stocks + Performance bonus
+ Health benefits ...

India - Pune / Hyd / Bangalore

Senior Architect / Principal Engineer @ google

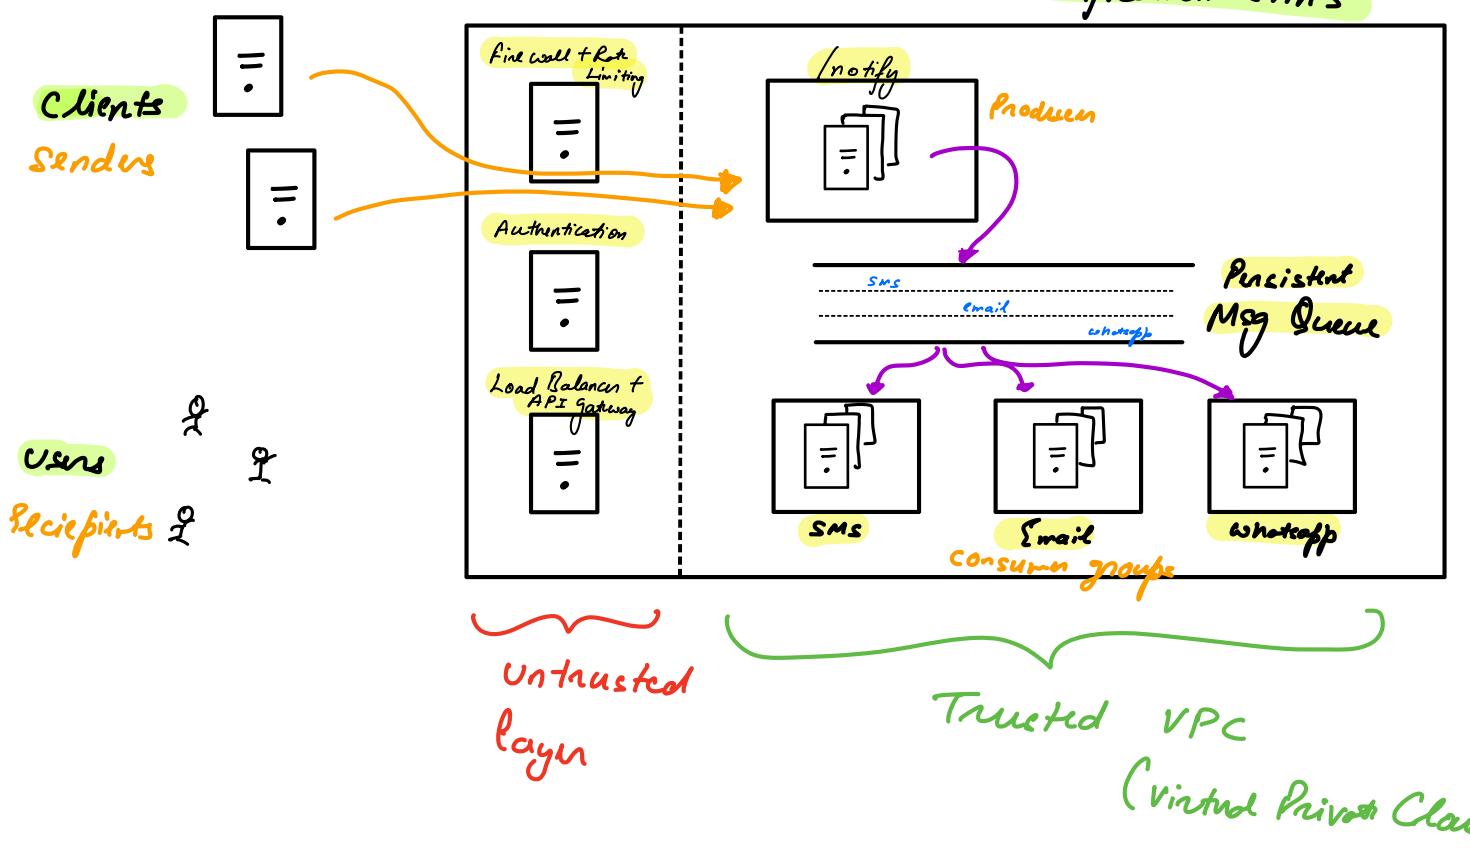
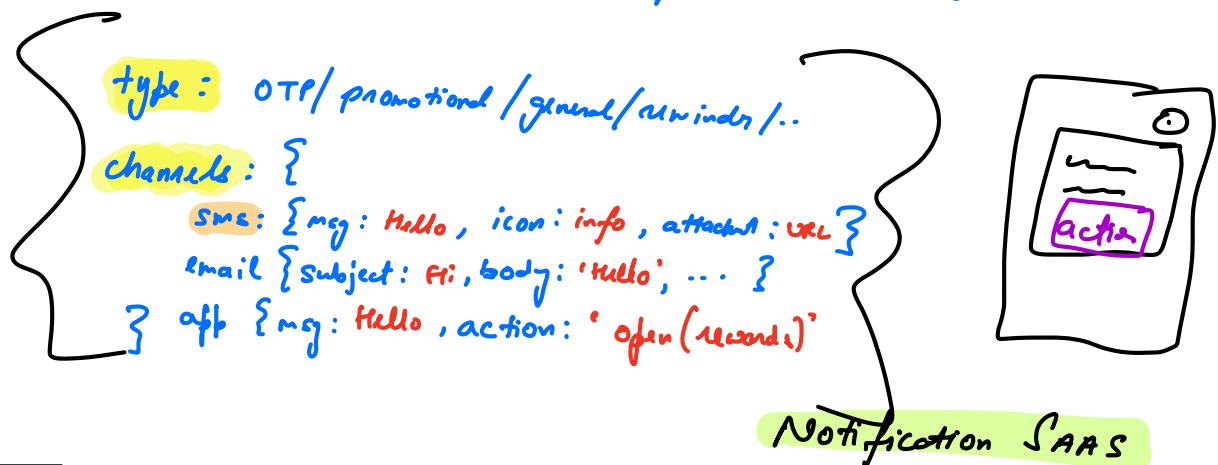
1 Cr - 3 Cr base

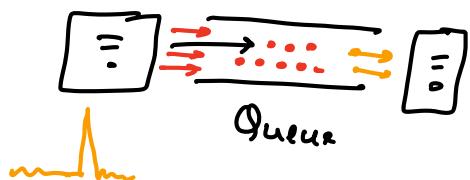
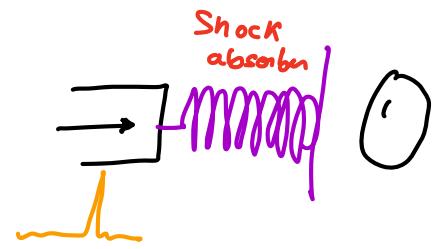
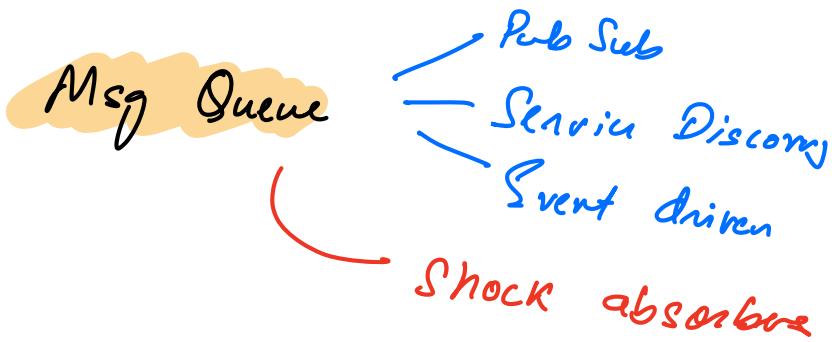
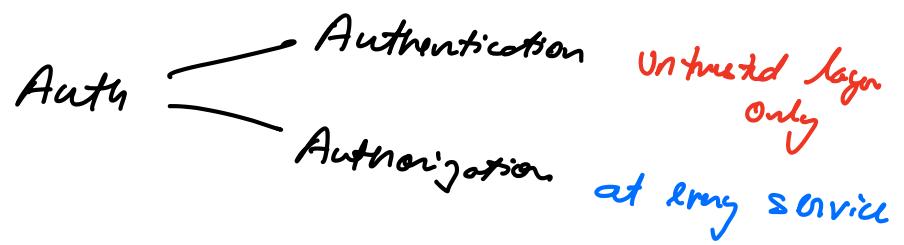
8:57 → 9:15

Break

System Architecture

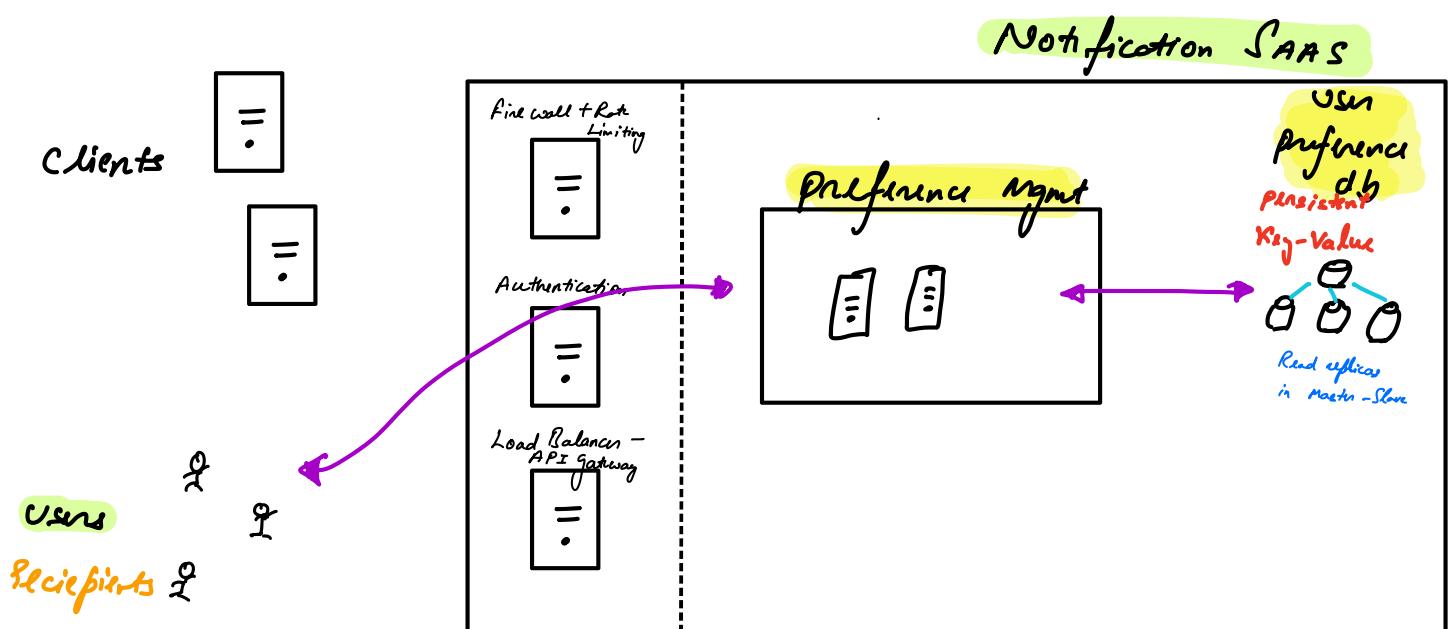
- ① clients should be able to send notifications
notify (api-key / user-id / content)
json
- ② system should support multiple channels





③ Users should be able to customize their notification preferences

System must respect the user settings



Ideal DB for user-preferences

	Strengths	Weaknesses
SQL PostgreSQL / MySQL / DB2 / RDS / ..	ACID Relations Normalization Schema Complex data model	Slow reads & writes don't scale well manual sharding

Key-Value Redis / Memcache / DynamoDB	Simple fast read optimized ✓	No relations Complex data modeling is not possible No search Single queries No transaction
---	------------------------------------	--

document Mongo / firebase / Elasticsearch / Couch / Cassandra / ...	Nested Semi-structured or schemaless data Search Powerful index & queries	No global index Slow writes No joins & relations No ACID transaction
--	--	---

wide column Cassandra / scylla / Hbase / BigTable	aggregate & analytics v. fast writes Pagination by time	No relations No ACID Modifying wide rows is slow
---	---	--

File Store / vector db / graph db - -

$$\cancel{2 \text{B user}} * 1 \text{KB preference data} = 2 \text{B} + 1 \text{KB data}$$

$$= 2 \text{TB}$$

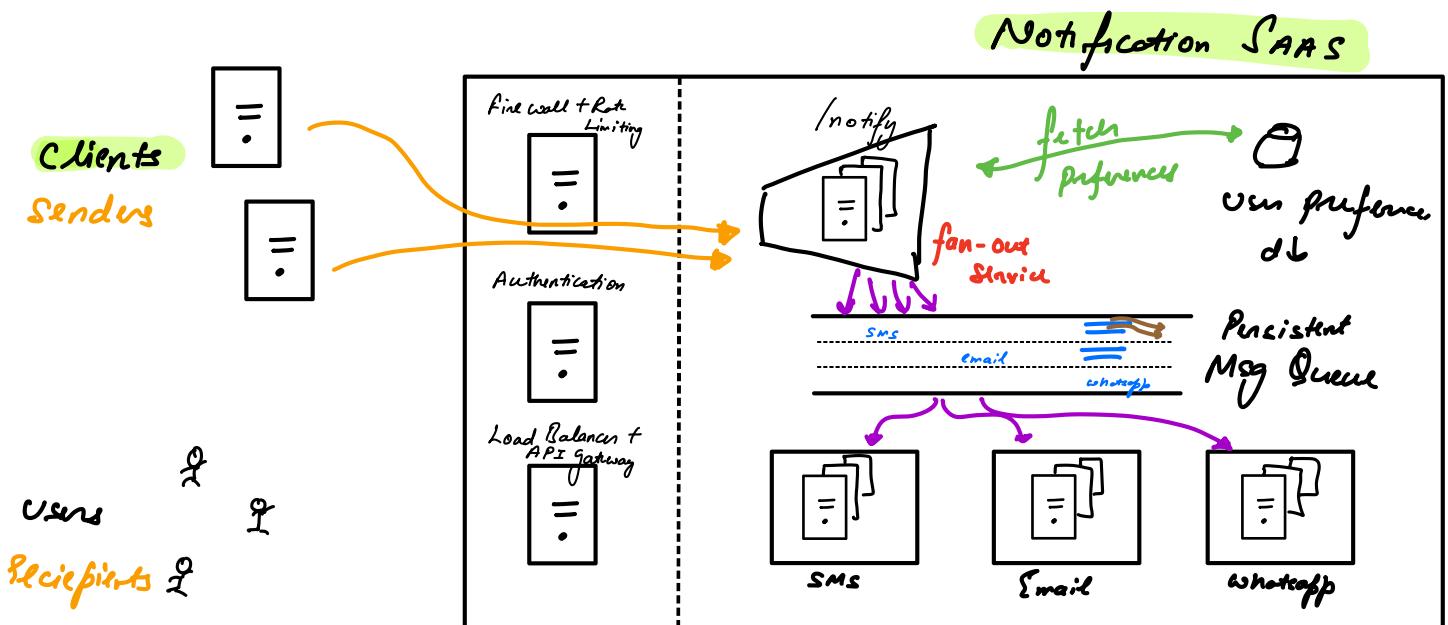
- Can be stored on a single server
- Read? v. High (10M/s)
- Write? low

Too much data Shard

Too many reads Cache
 Too many reads Read Replicas => improves reads
 but worsens writes

Too many writes Shard
 Too many writes Sample/Batch ∵ you have write more than once

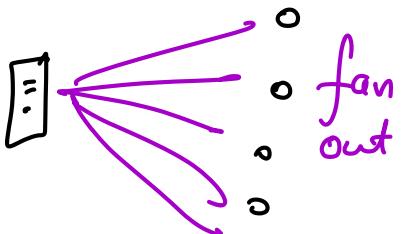
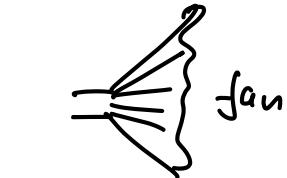
Prevent data loss Replicate



Fan-out notification service

for each notification req.

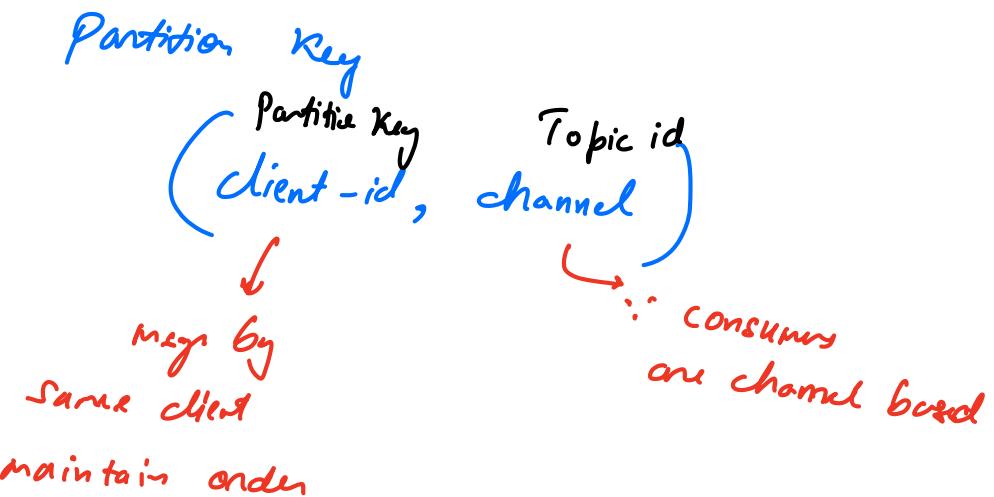
- Read the channels from request payload
- get user-id → fetch preferences from user-pref db
- filter out channels based on user pref
- broadcast messages (one for each channel) to msg Q



To separate user-db from pref management service & notification service *use CQRS*

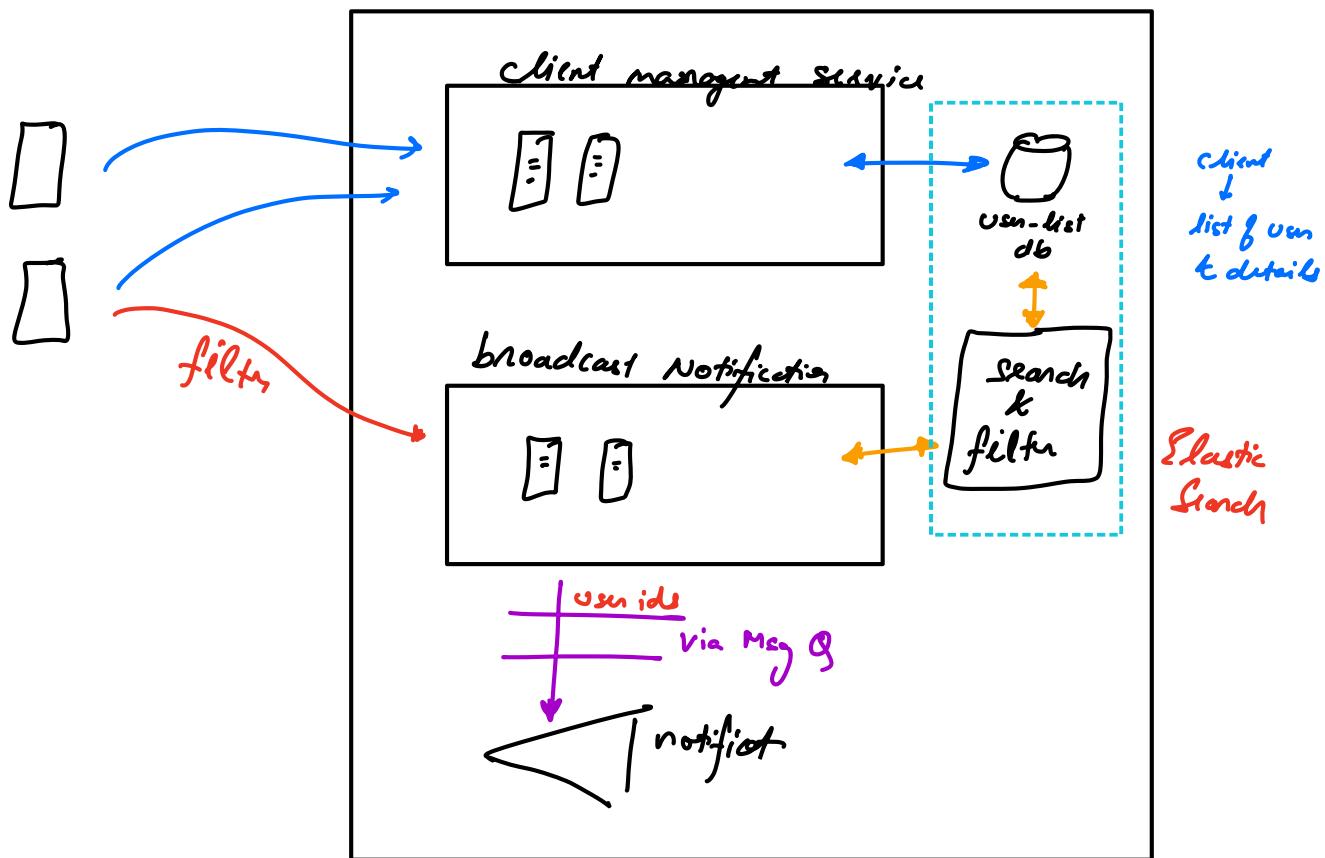
Message Queue

- ① Persistent — store notifications on disk (temporarily)
to ensure notification delivery.
- ② support offsets to ensure single delivery
- ③ Scattered → partition



⑦

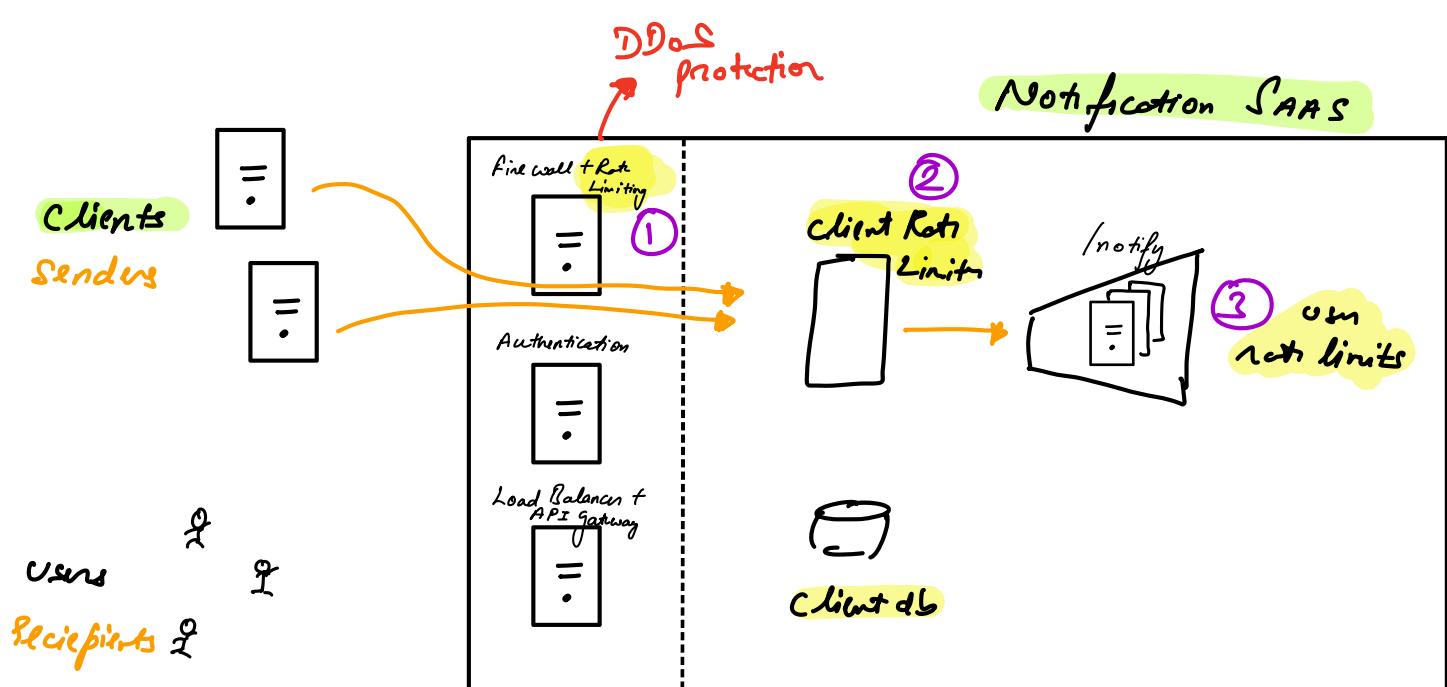
Clients should be able to broadcast msgs



⑥

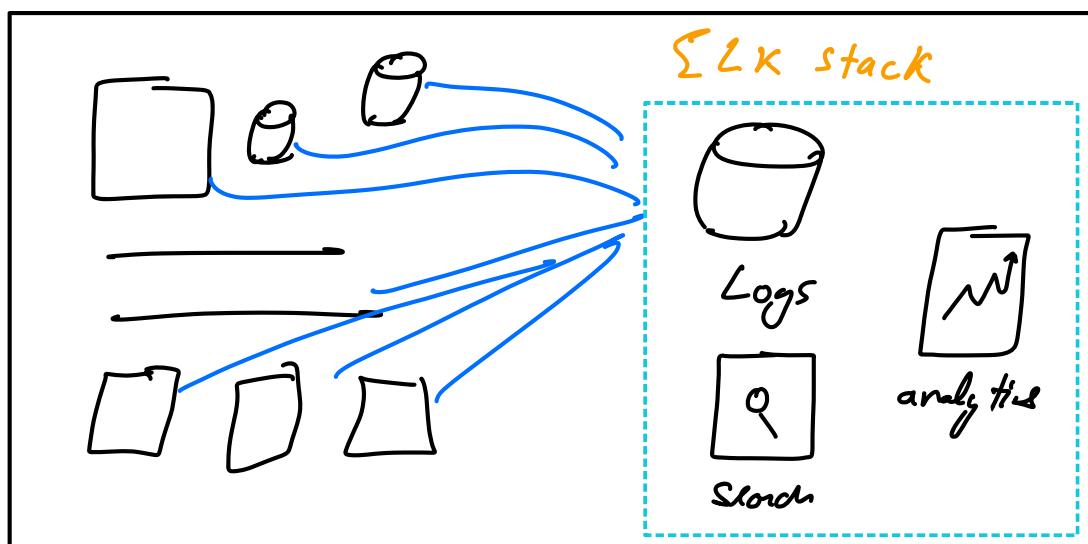
Rate limiting

- generic DDoS protection untrusted layer
- client
don't exceed purchased Quota
- User
don't spam user



⑦ Observability

- ① aggregate logs from all servers
- ② Search on logs
- ③ analytics (alerts / dashboards / monitoring)



Elastic Search

Log Stack

Kibana

Design Problems are open ended

- ① Rate limiting algos
- ② Event Driven archit
- ③ Celebrity Problem.

Elon Musk → 200M followers
M1 Rocket → 300M subscribers.

- ① Separate infra for celebrities

Twitter.

- ② Bypass the issue !!

Subscribe → Notification ✓

 Bell icon → notification ?